









# Deep Learning for Computer Vision

Dr. Konda Reddy Mopuri  
Mehta Family School of Data Science and Artificial Intelligence  
IIT Guwahati  
Aug-Dec 2022







# So far in the course

- Scoring function, loss function, gradient descent









# So far in the course

- Scoring function, loss function, gradient descent
- Artificial Neurons and Multi-Layered Perceptron



# So far in the course

- Scoring function, loss function, gradient descent
- Artificial Neurons and Multi-Layered Perceptron
- Backpropagation



# CNNs



- Neurons are similar to that of MLP

- Neurons are similar to that of MLP
  - Perform a linear (dot product) operation and have a nonlinearity

# CNNs



- Neurons are similar to that of MLP
  - Perform a linear (dot product) operation and have a nonlinearity
- Architecture will have a differentiable loss function, backpropagation is used

- Neurons are similar to that of MLP
  - Perform a linear (dot product) operation and have a nonlinearity
- Architecture will have a differentiable loss function, backpropagation is used
- Same tips and tricks apply




- Neurons are similar to that of MLP
  - Perform a linear (dot product) operation and have a nonlinearity
- Architecture will have a differentiable loss function, backpropagation is used
- Same tips and tricks apply
- So, what changes?

# An MLP




- Input is a vector



# An MLP

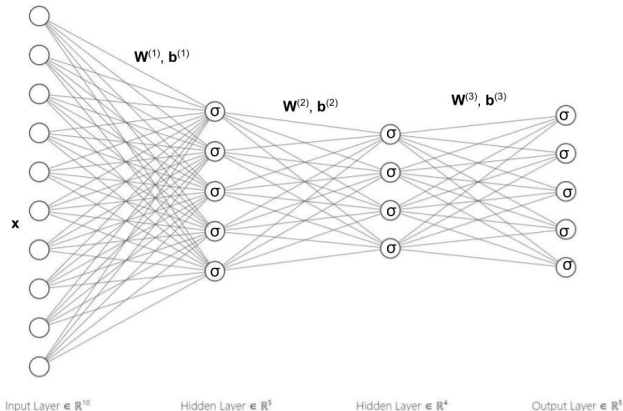


- Input is a vector
- Series of densely connected hidden layers



# An MLP

- Input is a vector
- Series of densely connected hidden layers
- Neurons in each layer are independent



# An MLP for processing an image



- Say, we want to process a  $200 \times 200$  RGB image

# An MLP for processing an image



- Say, we want to process a  $200 \times 200$  RGB image
- Vectorizing leads to  $200 \times 200 \times 3 \rightarrow 120K$  neurons in the input layer

# An MLP for processing an image



- Say, we want to process a  $200 \times 200$  RGB image
- Vectorizing leads to  $200 \times 200 \times 3 \rightarrow 120K$  neurons in the input layer
- A hidden layer of same size leads to  $\approx 1.44e^{10}$  weights  $\rightarrow \approx 58GB$

# An MLP for processing an image



- Say, we want to process a  $200 \times 200$  RGB image
- Vectorizing leads to  $200 \times 200 \times 3 \rightarrow 120K$  neurons in the input layer
- A hidden layer of same size leads to  $\approx 1.44e^{10}$  weights  $\rightarrow \approx 58GB$
- Full connectivity blows the number of weights  $\rightarrow$  hardware limits, overfitting, etc.



# An MLP for processing an image



- Say, we want to process a  $200 \times 200$  RGB image
- Vectorizing leads to  $200 \times 200 \times 3 \rightarrow 120K$  neurons in the input layer
- A hidden layer of same size leads to  $\approx 1.44e^{10}$  weights  $\rightarrow \approx 58GB$
- Full connectivity blows the number of weights  $\rightarrow$  hardware limits, overfitting, etc.
- Flattening removes the structure

# Large Signals



- Have invariance in translation

# Large Signals



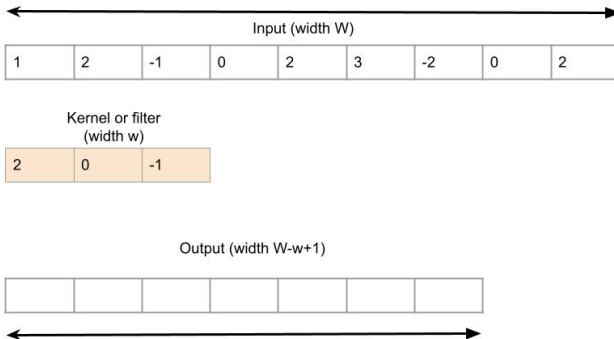
- Have invariance in translation
- Features may occur at different locations in the signal

# Large Signals

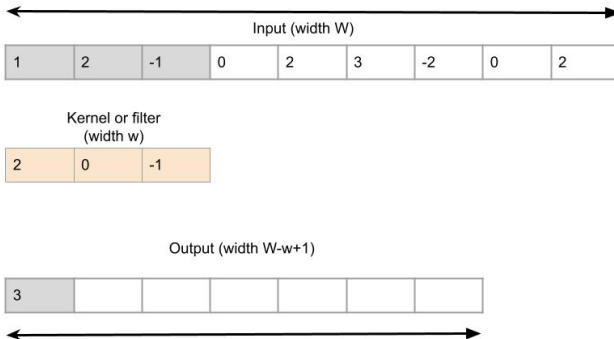


- Have invariance in translation
- Features may occur at different locations in the signal
- **Convolution** incorporates this idea: Applies same linear operation at all the locations and preserves the structure

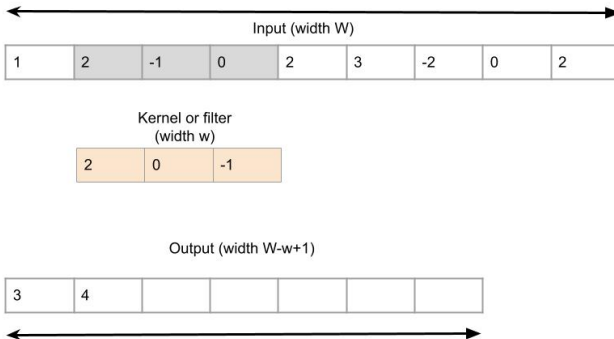
# Convolution



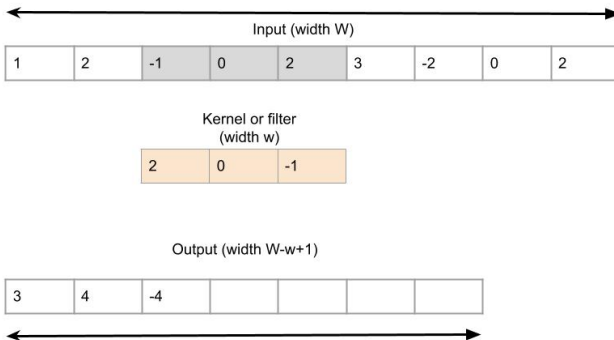
# Convolution



# Convolution

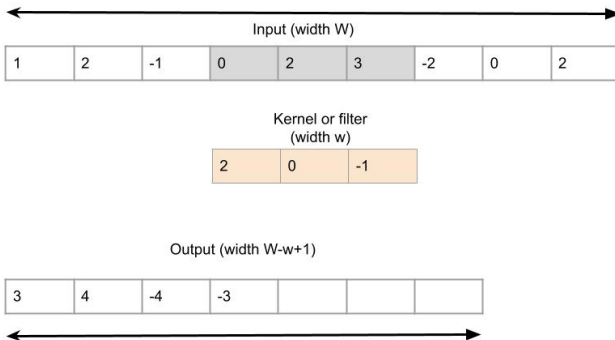


# Convolution







# Convolution




# Convolution



# Convolution



# Convolution



# Convolution



- Preserves the structure

# Convolution



- Preserves the structure
  - if the i/p is a 2D tensor  $\rightarrow$  o/p is also a 2D tensor

# Convolution



- Preserves the structure
  - if the i/p is a 2D tensor  $\rightarrow$  o/p is also a 2D tensor
  - There exist a relation between the locations of i/p and o/p values

# Convolution



- Let  $\mathbf{x} = (x_1, x_2, \dots, x_W)$  is the input,  $\mathbf{k} = (k_1, k_2, \dots, k_w)$  is the kernel



# Convolution



- Let  $\mathbf{x} = (x_1, x_2, \dots, x_W)$  is the input,  $\mathbf{k} = (k_1, k_2, \dots, k_w)$  is the kernel
- The result  $(x \circledast k)$  of convolving  $\mathbf{x}$  with  $\mathbf{k}$  will be a 1D tensor of size  $W - w + 1$

$$\begin{aligned}(x \circledast k)_i &= \sum_{j=1}^w x_{i-1+j} k_j \\ &= (x_i, \dots, x_{i+w-1}) \cdot \mathbf{k}\end{aligned}$$

# Convolution

- Powerful feature extractor



# Convolution



- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input


# Convolution



- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input



$$(0, 0, 0, 1, 2, 3, 4, 4, 4, 4) \otimes (-1, 1) = (0, 0, 1, 1, 1, 1, 0, 0, 0)$$




# Convolution



- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input



$$(0, 0, 1, 1, 0, 0.1, 0.2, 1, 1, 1, 0) \circledast (1, 1) = (0, 1, 2, 1, 0.1, 0.3, 1.2, 2, 2, 1)$$



# Convolution



- Naturally generalizes to multiple dimensions

# Convolution



- Naturally generalizes to multiple dimensions
- In their most usual form, CNNs process 3D tensors of size  $C \times H \times W$  with kernels of size  $C \times h \times w$  and result in 2D tensors of size  $H - h + 1 \times W - w + 1$

# Convolution





- Naturally generalizes to multiple dimensions
- In their most usual form, CNNs process 3D tensors of size  $C \times H \times W$  with kernels of size  $C \times h \times w$  and result in 2D tensors of size  $H - h + 1 \times W - w + 1$
- Note that we generally refer to these inputs as 2D signal (despite having  $C$  channels), because, they are referenced as vectors indexed by 2d locations without structure in the channel dimension



# 2D Convolution




input




kernel


# 2D Convolution




# 2D Convolution




# 2D Convolution




# 2D Convolution




# 2D Convolution



# 2D Convolution




# 2D Convolution







# 2D Convolution



# 2D Convolution



# 2D Convolution




# 2D Convolution



- Kernel is not convolved in the channel dimension


# 2D Convolution

- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size  $C \times h \times w$  and results in output of size  $D \times 1 \times 1$



# 2D Convolution


- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size  $C \times h \times w$  and results in output of size  $D \times 1 \times 1$



- Same affine function is applied on all such blocks in the input

# 2D Convolution

- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size  $C \times h \times w$  and results in output of size  $D \times 1 \times 1$



- Same affine function is applied on all such blocks in the input

# Convolution



- Preserves the input structure



# Convolution



- Preserves the input structure
  - 1D signal outputs 1D signal, 2D signal outputs 2D signal

# Convolution




- Preserves the input structure
  - 1D signal outputs 1D signal, 2D signal outputs 2D signal
  - Adjacent components in o/p are influenced by adjacent parts in the i/p

# Convolution



- Preserves the input structure
  - 1D signal outputs 1D signal, 2D signal outputs 2D signal
  - Adjacent components in o/p are influenced by adjacent parts in the i/p
- If the channel dimension has a metric meaning (e.g. time) 3D convolution can be employed (e.g. frames in a video)

# Terminology in Convolution



# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`

# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is  $D \times C \times h \times w$  dimensional kernels

# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is  $D \times C \times h \times w$  dimensional kernels
- `bias`  $D$  dimensional

# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is  $D \times C \times h \times w$  dimensional kernels
- `bias`  $D$  dimensional
- `input` is  $N \times C \times H \times W$  dimensional signal



# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is  $D \times C \times h \times w$  dimensional kernels
- `bias`  $D$  dimensional
- `input` is  $N \times C \times H \times W$  dimensional signal
- Output is  $N \times D \times (H - h + 1) \times (W - w + 1)$  tensor

# Convolution function in PyTorch



- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is  $D \times C \times h \times w$  dimensional kernels
- `bias`  $D$  dimensional
- `input` is  $N \times C \times H \times W$  dimensional signal
- Output is  $N \times D \times (H - h + 1) \times (W - w + 1)$  tensor
- Autograd compliant

# Convolution function in PyTorch

```
input = torch.empty(128, 3, 20, 20).normal_()
weight = torch.empty(5, 3, 5, 5).normal_()
bias = torch.empty(5).normal_()
output = F.conv2d(input, weight, bias)
output.size()
torch.Size([128, 5, 16, 16])
```

# Look/Access the filters



```
weight[0,0]  
tensor([[ -0.6974,  0.1342, -0.2632, -0.4672,  0.1827],  
        [ -0.1184, -0.2164,  0.2772, -0.1099,  0.0103],  
        [ -0.8272,  0.3580,  0.2398, -0.5795, -0.9472],  
        [ -1.1734, -0.1019,  0.7394,  0.3342,  0.1699],  
        [  1.9271,  0.1250,  0.4222,  0.2014,  1.1100]])
```

# Conv layer in PyTorch



- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`

# Conv layer in PyTorch



- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair  $(h, w)$  or a single value  $k$  interpreted as  $(k, k)$ .

# Conv layer in PyTorch



- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair  $(h, w)$  or a single value  $k$  interpreted as  $(k, k)$ .
- Encloses the convolution as a module

# Conv layer in PyTorch



- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair  $(h, w)$  or a single value  $k$  interpreted as  $(k, k)$ .
- Encloses the convolution as a module
- Initializes the kernel parameters and biases as random





# Conv layer in PyTorch

```
f = nn.Conv2d(in_channels = 3, out_channels = 5,  
kernel_size = (2, 3))  
for n, p in f.named_parameters():  
    ...print(n, p.size())  
  
>>weight torch.Size([5, 3, 2, 3])  
>>bias torch.Size([5])
```

# Conv layer in PyTorch

```
f = nn.Conv2d(in_channels = 3, out_channels = 5,  
kernel_size = (2, 3))  
for n, p in f.named_parameters():  
    ...print(n, p.size())  
  
>>weight torch.Size([5, 3, 2, 3])  
>>bias torch.Size([5])  
  
input = torch.empty(128, 3, 28, 28).normal_()  
output = f(input)  
output.size()  
>>torch.Size([128, 5, 27, 26])
```

# Padding in Convolution



- Adds zeros around the input

# Padding in Convolution




- Adds zeros around the input
- Takes care of size reduction after convolution

# Padding in Convolution

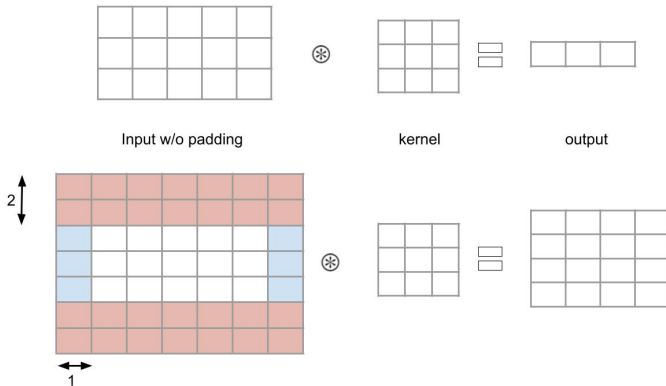


- Adds zeros around the input
- Takes care of size reduction after convolution
- Instead of zeros, one may pad with signal values at the edges

# Padding in Convolution



# Padding in Convolution



# Stride in Convolution



- Specifies the step size taken while performing convolution




# Stride in Convolution



- Specifies the step size taken while performing convolution
- Default value is 1, i.e., move the kernel across the signal densely (without skipping)

# Padding and Stride in Convolution



# Dilation in Convolution




- Manipulates the size of the kernel via expanding its size without adding weights.

# Dilation in Convolution

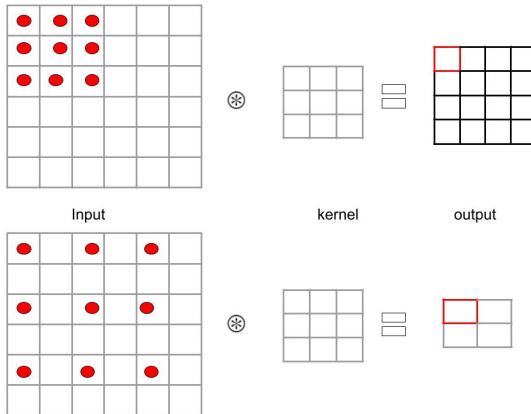


- Manipulates the size of the kernel via expanding its size without adding weights.
- In other words, it inserts 0s in between the kernel values

# Without Dilation



# Dilation (2, 2)



# Dilation



- Expands the kernel by adding rows and columns of zeros

# Dilation



- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed



# Dilation



- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse

# Dilation



- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse
- Dilation increases the receptive field

# Dilation



- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse
- Dilation increases the receptive field
- It is referred to as 'atrous' convolution

# Pooling

# Pooling



- Groups multiple activations and replaces by a representative one

# Pooling



- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride

# Pooling



- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride
- Also called sub-sampling layer

# Pooling



- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride
- Also called sub-sampling layer
- Generally found between two convolution layers (and parameter free)




# Max Pooling



- Standard in CNNs

# Max Pooling


- Standard in CNNs
- Computes maximum value over a non-overlapping blocks in the input



# Average Pooling



- Computes the average of the receptive field



# Pooling in 2D

- Same as 1D, but the receptive field is 2D and non-overlapping




Figure credits: Preston Hoang and Quora



# Pooling in 2D



- Contrary to Convolution, Pooling applies channel wise

# Pooling in 2D

- Contrary to Convolution, Pooling applies channel wise
- No reduction in number of channels, only spatial size reduction



# Pooling provides weak invariance



- Operation is invariant to any permutation within the block

# Pooling provides weak invariance



- Operation is invariant to any permutation within the block
- Withstands deformations caused by local translations



# Max\_Pooling PyTorch



```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately

# Max\_Pooling PyTorch



```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- input is  $N \times C \times H \times W$  tensor

# Max\_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- $\text{input}$  is  $N \times C \times H \times W$  tensor
- $\text{kernel\_size}$  is  $(h, w)$  or  $k$

# Max\_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- input is  $N \times C \times H \times W$  tensor
- kernel\_size is  $(h, w)$  or  $k$
- Result would be a tensor of size  $N \times C \times \lfloor H/h \rfloor \times \lfloor W/w \rfloor$

# Pooling in PyTorch



- Default stride is the kernel size (for convolution, it is 1)

# Pooling in PyTorch



- Default stride is the kernel size (for convolution, it is 1)
- But, it can be modulated if required

# Pooling in PyTorch



- Default stride is the kernel size (for convolution, it is 1)
- But, it can be modulated if required
- Default padding is zero

# Pooling Layer in PyTorch



```
class torch.nn.MaxPool2d(kernel_size, stride=None,  
padding=0, dilation=1, return_indices=False,  
ceil_mode=False)
```



# Putting it all together

# Architecture of a simple CNN





Figure credits: Adit Deshpande


# Architecture of a simple CNN



- Initially Conv layer with nonlinearity

Figure credits: Adit Deshpande


# Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers

Figure credits: Adit Deshpande


# Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers
- Have Pooling layers in between Conv layers → reduce the feature map size sufficiently

Figure credits: Adit Deshpande


# Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers
- Have Pooling layers in between Conv layers → reduce the feature map size sufficiently
- Vectorize and fully connected layers

Figure credits: Adit Deshpande


# Architecture of a simple CNN



INPUT  $\rightarrow$   $[[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL}] * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow \text{FC}$

Figure credits: Adit Deshpande

# Architecture of a simple CNN




INPUT  $\rightarrow$   $[[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL}] * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow$   
FC

Figure credits: Adit Deshpande



# Architecture of a simple CNN



INPUT  $\rightarrow$   $[[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL}] * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow$   
FC

Figure credits: Adit Deshpande

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>			

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$		

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1) = 832$	

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1) = 832$	$32 \cdot 24^2 \cdot 5^2 = 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>			

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ nn.Conv2d(1, 32, kernel_size=5)	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2 = 460800$
$32 \times 24 \times 24$ F.max_pool2d(., kernel_size=3)	$32 \times 8 \times 8$	0	0

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ nn.Conv2d(1, 32, kernel_size=5)	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2 = 460800$
$32 \times 24 \times 24$ F.max_pool2d(., kernel_size=3)	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / F.relu(.)	$32 \times 8 \times 8$	0	0

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>			



# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$		

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	$= 460800$ 0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	$= 460800$ 0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200		

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	$= 460800$ 0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200	$200(256+1)=51400$	$200.256=51200$

# Case study: LeNet-like architecture



input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1) = 832$	$32.24^2.5^2$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	$= 460800$ 0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.4^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200	0 $200(256+1)=51400$	0 $200.256=51200$
$200$ / <code>F.relu(.)</code>	200	0	0
200 <code>nn.Linear(200,10)</code>	10	0 $10(200+1)=2010$	0 $10.200=2000$



# Recent architectures are more sophisticated



- Note that LeNet is a classical architecture and does not reflect the recent CNNs in complexity

# Recent architectures are more sophisticated



- Note that LeNet is a classical architecture and does not reflect the recent CNNs in complexity
- Recent CNN architectures are far more sophisticated [Contents of the next lecture(s)]
  - More depth
  - Regularizers to handle the depth