



# Deep Learning for Computer Vision

Dr. Konda Reddy Mopuri  
Mehta Family School of Data Science and Artificial Intelligence  
IIT Guwahati  
Aug-Dec 2022

# So far in the class..



- Brief introduction to ML

# So far in the class..



- Brief introduction to ML
- Artificial neuron models, Perceptron

# So far in the class..



- Brief introduction to ML
- Artificial neuron models, Perceptron
- MLP, CNNs and different families of architecture

# So far in the class..



- Brief introduction to ML
- Artificial neuron models, Perceptron
- MLP, CNNs and different families of architecture
- (today) Some of the important training aspects of CNNs

# Data preprocessing for Computer vision



- Mean subtraction (e.g. AlexNet:  $32 \times 32 \times 3$ , VGG:  $1 \times 1 \times 3$ )

# Data preprocessing for Computer vision



- Mean subtraction (e.g. AlexNet:  $32 \times 32 \times 3$ , VGG:  $1 \times 1 \times 3$ )
- Mean subtraction and division by standard deviation per channel (e.g. ResNet)

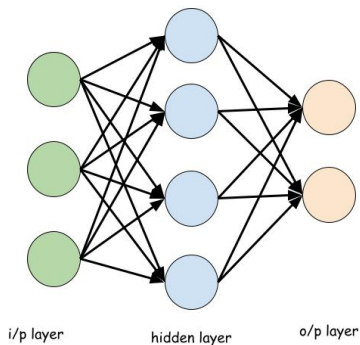
# Data preprocessing for Computer vision



- Mean subtraction (e.g. AlexNet:  $32 \times 32 \times 3$ , VGG:  $1 \times 1 \times 3$ )
- Mean subtraction and division by standard deviation per channel (e.g. ResNet)
- PCA or whitening are not common




# Weight Initialization



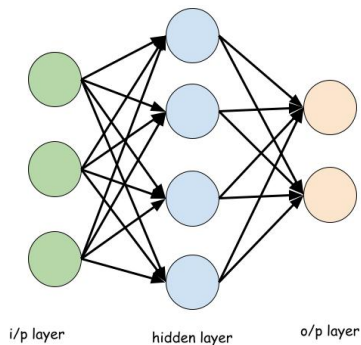
- What if all the parameters are initialized to zero?

# Weight Initialization



- What if all the parameters are initialized to zero?
- Or, a different constant?

# Weight Initialization



- What if all the parameters are initialized to zero?
- Or, a different constant?
- Leads to a failure mode (often known as the 'symmetry' problem)

# Weight Initialization

- How about randomly initializing?

$$W = 0.001 * \text{np.random.randn}(d_l, d_{l-1})$$

---

Figure credits: Dr Justin Johnson, U Michigan

# Weight Initialization

- How about randomly initializing?  
 $W = 0.001 * \text{np.random.randn}(d_l, d_{l-1})$
- Okay for the shallow nets

---

Figure credits: Dr Justin Johnson, U Michigan

# Weight Initialization

- How about randomly initializing?  
 $W = 0.001 * \text{np.random.randn}(d_l, d_{l-1})$
- Okay for the shallow nets
- However, the dynamic range of the activations at later layers goes on shrinking  $\rightarrow$  activations tend to zero at deeper layers (e.g. 6 layer MLP with a tanh nonlinearity)

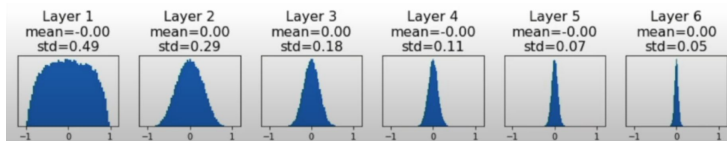



Figure credits: Dr Justin Johnson, U Michigan

# Weight Initialization

- How about randomly initializing?  
 $W = 0.001 * \text{np.random.randn}(d_l, d_{l-1})$
- Okay for the shallow nets
- However, the dynamic range of the activations at later layers goes on shrinking  $\rightarrow$  activations tend to zero at deeper layers (e.g. 6 layer MLP with a tanh nonlinearity)



- All zero gradients, no learning!

Figure credits: Dr Justin Johnson, U Michigan

# Xavier Initialization



- $W = 0.001 * \text{np.random.randn}(d_l, d_{l-1}) / \text{np.sqrt}(d_{l-1})$

---

Figure credits: Dr Justin Johnson, U Michigan



# Xavier Initialization



- $W = 0.001 * \text{np.random.randn}(d_l, d_{l-1}) / \text{np.sqrt}(d_{l-1})$

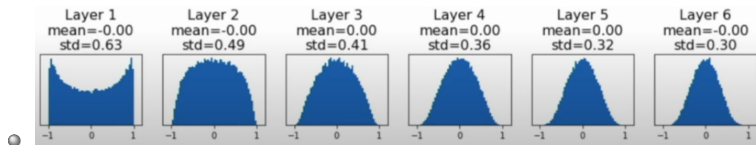


Figure credits: Dr Justin Johnson, U Michigan

# Xavier Initialization



- We prefer the o/p to have similar variance as the input

# Xavier Initialization



- We prefer the o/p to have similar variance as the input
- Consider a single layer,  $y = Wx$ , i.e.  $y_i = \sum_{j=1}^{d_{l-1}} x_j \cdot w_j$

# Xavier Initialization



- We prefer the o/p to have similar variance as the input
- Consider a single layer,  $y = Wx$ , i.e.  $y_i = \sum_{j=1}^{d_{l-1}} x_j \cdot w_j$
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i \cdot w_i)$  (Assuming  $w_i$  and  $x_i$  are i.i.d)

# Xavier Initialization



- We prefer the o/p to have similar variance as the input
- Consider a single layer,  $y = Wx$ , i.e.  $y_i = \sum_{j=1}^{d_{l-1}} x_j \cdot w_j$
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i \cdot w_i)$  (Assuming  $w_i$  and  $x_i$  are i.i.d)
- $\text{var}(y_i) = d_{l-1} \cdot \left( E(x_i^2) \cdot E(w_i^2) - E(x_i)^2 \cdot E(w_i)^2 \right)$  (Assuming  $x$  and  $w$  are independent)

# Xavier Initialization



- We prefer the o/p to have similar variance as the input
- Consider a single layer,  $y = Wx$ , i.e.  $y_i = \sum_{j=1}^{d_{l-1}} x_j \cdot w_j$
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i \cdot w_i)$  (Assuming  $w_i$  and  $x_i$  are i.i.d)
- $\text{var}(y_i) = d_{l-1} \cdot \left( E(x_i^2) \cdot E(w_i^2) - E(x_i)^2 \cdot E(w_i)^2 \right)$  (Assuming  $x$  and  $w$  are independent)
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i) \cdot \text{var}(w_i)$  Assuming ( $x_i$  and  $w_i$  are zero-mean)

# Xavier Initialization



- We prefer the o/p to have similar variance as the input
- Consider a single layer,  $y = Wx$ , i.e.  $y_i = \sum_{j=1}^{d_{l-1}} x_j \cdot w_j$
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i \cdot w_i)$  (Assuming  $w_i$  and  $x_i$  are i.i.d)
- $\text{var}(y_i) = d_{l-1} \cdot \left( E(x_i^2) \cdot E(w_i^2) - E(x_i)^2 \cdot E(w_i)^2 \right)$  (Assuming  $x$  and  $w$  are independent)
- $\text{var}(y_i) = d_{l-1} \cdot \text{var}(x_i) \cdot \text{var}(w_i)$  Assuming ( $x_i$  and  $w_i$  are zero-mean)
- $\rightarrow \text{var}(w_i) = \frac{1}{d_{l-1}}$

# Weight Initialization with ReLU activations



- Kaiming He or MSRA initialization

---

Figure credits: Dr Justin Johnson



# Weight Initialization with ReLU activations



- Kaiming He or MSRA initialization
- $\text{std} = \sqrt{2/d_{l-1}}$





Figure credits: Dr Justin Johnson


# Weight Initialization: Residual Networks



- MSRA initialization:  
 $\text{Var}(F(x)+x) > \text{Var}(x)$

Figure credits: Dr. Justin Johnson


# Weight Initialization: Residual Networks



- MSRA initialization:  
 $\text{Var}(F(x)+x) > \text{Var}(x)$
- Variance grows!

Figure credits: Dr. Justin Johnson

# Weight Initialization: Residual Networks



- MSRA initialization:  
 $\text{Var}(F(x)+x) > \text{Var}(x)$
- Variance grows!
- Solution: Initialize the first Conv layer with MSRA, and the second one with zero  $\rightarrow$   
 $\text{Var}(x+F(x)) = \text{Var}(x)$

Figure credits: Dr. Justin Johnson

# Deep Regularization



- ① Most of the regularization techniques for deep learning are based on regularizing estimators

# Deep Regularization



- ① Most of the regularization techniques for deep learning are based on regularizing estimators
- ② Trade increased bias for decreased variance

# Deep Regularization



- ① An overly complex model family need not include the target function

# Deep Regularization



- ① An overly complex model family need not include the target function
- ② In practice we almost never have access to the true data generating process, and which is almost certainly outside the model family



# Deep Regularization



- ① Most often the best-fitting model is a large model that has been appropriately regularized

# Deep Regularization



- Parameter Norm penalties ( $l_2$ ,  $l_1$ , etc.)
- Dataset Augmentation
- Noise Robustness
- Semi-Supervised Learning
- Multi-Task Learning (Parameter sharing)
- Sparse Representation
- Dropout
- etc.

# Parameter Norm Penalties



- ① For neural networks, typically only the weights of the affine transformations are regularized leaving the biases unregularized

# Parameter Norm Penalties



- ① For neural networks, typically only the weights of the affine transformations are regularized leaving the biases unregularized
- ② Bias controls only a single variable as opposed to weight which connects two

# Parameter Norm Penalties



- ① For neural networks, typically only the weights of the affine transformations are regularized leaving the biases unregularized
- ② Bias controls only a single variable as opposed to weight which connects two
- ③ Regularizing biases may induce underfitting

# Parameter Norm Penalties



- ①  $L_2$  parameter regularization:  $\tilde{\mathcal{J}} = \frac{\alpha}{2} w^T w + \mathcal{J}(w; X, y)$

# Parameter Norm Penalties



- ①  $L_2$  parameter regularization:  $\tilde{\mathcal{J}} = \frac{\alpha}{2} w^T w + \mathcal{J}(w; X, y)$
- ②  $L_1$  regularization:  $\tilde{\mathcal{J}} = \alpha |w|_1 + \mathcal{J}(w; X, y)$

# Parameter Norm Penalties



- ①  $L_2$  parameter regularization:  $\tilde{\mathcal{J}} = \frac{\alpha}{2} w^T w + \mathcal{J}(w; X, y)$
- ②  $L_1$  regularization:  $\tilde{\mathcal{J}} = \alpha |w|_1 + \mathcal{J}(w; X, y)$
- ③ Norm penalties induce different desired behaviors based on the exact penalty imposed



# Dataset Augmentation



- ① Bestway to make ML model generalize better is to train with more data

# Dataset Augmentation



- ① Bestway to make ML model generalize better is to train with more data
- ② In practice training data is limited

# Dataset Augmentation



- ① Bestway to make ML model generalize better is to train with more data
- ② In practice training data is limited
- ③ Create fake data and add it to the training data, called Dataset augmentation

# Dataset Augmentation



- ① Easier for classification

# Dataset Augmentation



- ① Easier for classification
- ② Difficult for density estimation task (unless we have solved the estimation problem)

# Dataset Augmentation



- ① Has been particularly effective for specific classification problems such as object recognition

# Dataset Augmentation



- ① Has been particularly effective for specific classification problems such as object recognition
- ② Operations such as translation by few pixels, rotating slightly, adding mild noise, etc. greatly improve generalization

# Dataset Augmentation



- ① Has been particularly effective for specific classification problems such as object recognition
- ② Operations such as translation by few pixels, rotating slightly, adding mild noise, etc. greatly improve generalization
- ③ Hand-designed augmentations in some domains can result in dramatic improvements
- ④ Should restrict to label preserving transformations



# Multi-Task Learning



- ① Improves generalization by collecting samples arising out of multiple tasks


# Multi-Task Learning



- ① Improves generalization by collecting samples arising out of multiple tasks
- ② Similar to additional data samples, multi-task samples also put more pressure on the parameters of the shared layers to be better

# Multi-Task Learning

- ① Improves generalization by collecting samples arising out of multiple tasks
- ② Similar to additional data samples, multi-task samples also put more pressure on the parameters of the shared layers to be better



# Dropout



- ① Key ideas and contributions in DL have been to engineer architectures for making them easier to train

# Dropout



- ① Key ideas and contributions in DL have been to engineer architectures for making them easier to train
- ② Dropout is one such ('deep') regularization technique (Srivastava et al. 2014)

# Dropout



- ① During the forward pass, some of the units are randomly 'zeroed' out (neurons are removed)




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Figure from Srivastava et al. 2014

# Dropout



- ① During the forward pass, some of the units are randomly 'zeroed' out (neurons are removed)
- ② Dropped units are randomly selected in each layer independent of others




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Figure from Srivastava et al. 2014

# Dropout



- ① During the forward pass, some of the units are randomly 'zeroed' out (neurons are removed)
- ② Dropped units are randomly selected in each layer independent of others
- ③ Resulting network has a different architecture




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Figure from Srivastava et al. 2014



# Dropout



- ① During the forward pass, some of the units are randomly 'zeroed' out (neurons are removed)
- ② Dropped units are randomly selected in each layer independent of others
- ③ Resulting network has a different architecture
- ④ Backpropagation happens through the remaining activations




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Figure from Srivastava et al. 2014

# Dropout: Interpretation



- ① Improves independence between the units (prevents co-adaptation of the units in the network)

# Dropout: Interpretation



- ① Improves independence between the units (prevents co-adaptation of the units in the network)
- ② Distributes the representation among all the units (forces the network to learn redundancy)

# Dropout



- ① We will decide on which units/layers to use dropout, and with what probability  $p$  units are dropped.

# Dropout



- ① We will decide on which units/layers to use dropout, and with what probability  $p$  units are dropped.
- ② For each sample, as many Bernoulli variables as units are sampled independently for dropping the units.

# Dropout: Another Interpretation



- ① Results in a large ensemble of networks (with shared parameters)

# Dropout: Another Interpretation



- ① Results in a large ensemble of networks (with shared parameters)
- ② Every possible binary mask results in a member of the ensemble

# Dropout: Another Interpretation



- ① Results in a large ensemble of networks (with shared parameters)
- ② Every possible binary mask results in a member of the ensemble
- ③ E.g. a dense layer with 10 units has  $2^{10}$  masks!



# Dropout: test time



- ① Which model from the ensemble to use?  
 $y = f(x, w, m)$  ( $m$  is the chosen binary mask)

# Dropout: test time

- ① Which model from the ensemble to use?  
 $y = f(x, w, m)$  ( $m$  is the chosen binary mask)
- ② How about taking the opinion of all the experts? → 'average out' and make the o/p deterministic

# Dropout: test time

- ① Which model from the ensemble to use?  
 $y = f(x, w, m)$  ( $m$  is the chosen binary mask)
- ② How about taking the opinion of all the experts? → 'average out' and make the o/p deterministic
- ③  $y = \mathbb{E}_m[f(x, w, m)] = \sum_m p(m) \cdot f(x, w, m)$

# Dropout: test time

- ① Which model from the ensemble to use?  
 $y = f(x, w, m)$  ( $m$  is the chosen binary mask)
- ② How about taking the opinion of all the experts? → 'average out' and make the o/p deterministic
- ③  $y = \mathbb{E}_m[f(x, w, m)] = \sum_m p(m) \cdot f(x, w, m)$
- ④ Leads to dropping no unit but multiply the activations with the probability of retaining

# Dropout: test time

- ① Which model from the ensemble to use?  
 $y = f(x, w, m)$  ( $m$  is the chosen binary mask)
- ② How about taking the opinion of all the experts? → 'average out' and make the o/p deterministic
- ③  $y = \mathbb{E}_m[f(x, w, m)] = \sum_m p(m) \cdot f(x, w, m)$
- ④ Leads to dropping no unit but multiply the activations with the probability of retaining
- ⑤ The standard variant uses the 'inverted dropout'. Multiplies activations by  $\frac{1}{(1-p)}$  during train and keeps the network untouched during test.

# Dropout



① Which layers to regularize with the Dropout?

# Dropout



- ① Which layers to regularize with the Dropout?
- ② More parameters are the dense layers → usually applied there

# Dropout



- ① Which layers to regularize with the Dropout?
- ② More parameters are the dense layers → usually applied there
- ③ Not much used after ResNets!



# Batch Normalization (BN)



- ① Gradient Descent converges faster with feature scaling ( $x \leftarrow \frac{x-\mu}{\sigma}$ )

# Batch Normalization (BN)



- ① Gradient Descent converges faster with feature scaling ( $x \leftarrow \frac{x-\mu}{\sigma}$ )
- ② Batch Normalization (BN) is a normalization method for intermediate layers of NNs  $\rightarrow$  performs whitening to the intermediate layer activations

# Batch Normalization (BN)

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1...m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

$\gamma$  and  $\beta$  are learn-able parameters

# Batch Normalization (BN)



- 1 Originally introduced to handle the internal covariate shift (ICS)

# Batch Normalization (BN)



- ① Originally introduced to handle the internal covariate shift (ICS)
- ② BN makes the activation of each neuron to be Gaussian distributed

# Batch Normalization (BN)




- ① Originally introduced to handle the internal covariate shift (ICS)
- ② BN makes the activation of each neuron to be Gaussian distributed
- ③ ICS is undesirable because the layers need to adapt to the new distribution of activations
- ④ With BN, it is reduced to new pair of parameters, but the distribution remains Gaussian

# Batch Normalization (BN)




- ① Mitigates interdependency between hidden layers during training



# Batch Normalization (BN)



- ① Mitigates interdependency between hidden layers during training




②  $\partial(a) = \partial(b) \cdot \partial(c) \cdot \partial(d) \cdot \partial(e)$



# Batch Normalization (BN)



- ① Mitigates interdependency between hidden layers during training




- ②  $\partial(a) = \partial(b) \cdot \partial(c) \cdot \partial(d) \cdot \partial(e)$
- ③ if we want to adjust the input distribution of a specific hidden unit, we need to consider the whole sequence of layers (w/o BN)

# Batch Normalization (BN)



- ① Mitigates interdependency between hidden layers during training



- ②  $\partial(a) = \partial(b) \cdot \partial(c) \cdot \partial(d) \cdot \partial(e)$
- ③ if we want to adjust the input distribution of a specific hidden unit, we need to consider the whole sequence of layers (w/o BN)
- ④ BN acts like a valve which holds back the flow, and allows its regulation using  $\beta$  and  $\gamma$

# Batch Normalization (BN)



- ① Reduces training time (less ICS)

# Batch Normalization (BN)



- ① Reduces training time (less ICS)
- ② Reduces the demand for additional regularizers (Batch statistics)

# Batch Normalization (BN)



- ① Reduces training time (less ICS)
- ② Reduces the demand for additional regularizers (Batch statistics)
- ③ Allows higher learning rates (less danger of vanishing/exploding gradients)

# Regularization: General idea



- 1 Add some randomness during the training

# Regularization: General idea



- ① Add some randomness during the training
- ② Have a mechanism for marginalizing while testing

# Regularization: General idea

- ① Add some randomness during the training
- ② Have a mechanism for marginalizing while testing
- ③ Some of the instances

Dropout

Batch Normalization

Data Augmentation

Drop Connect (drop weights instead)

Fractional MaxPooling

Stochastic Depth

Mixup

Cutout

CutMix, etc.