



Deep Learning for Computer Vision

Dr. Konda Reddy Mopuri
Mehta Family School of Data Science and Artificial Intelligence
IIT Guwahati
Aug-Dec 2022

Autoencoders



- ① Designed to reproduce input, especially reproduce the input from a learned encoding

Autoencoders



- ① Designed to reproduce input, especially reproduce the input from a learned encoding
- ② We attempted to project the data into the latent space and model it via a probability distribution

Autoencoders




- ① Designed to reproduce input, especially reproduce the input from a learned encoding
- ② We attempted to project the data into the latent space and model it via a probability distribution
- ③ This wasn't satisfying

Variational Autoencoders




- ① Key idea is to make both Encoder and Decoder stochastic




Variational Autoencoders

- ① Key idea is to make both Encoder and Decoder stochastic
- ② Latent variable z is drawn from a probability distribution for the given input x



Variational Autoencoders

- ① Key idea is to make both Encoder and Decoder stochastic
- ② Latent variable z is drawn from a probability distribution for the given input x
- ③ Also, the reconstruction is chosen probabilistically from the sampled z



VAE Encoder



- ① Takes input and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)

VAE Encoder



- ① Takes input and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)
- ② We can sample this to get random values of the latent variable z

VAE Encoder



- ① Takes input and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)
- ② We can sample this to get random values of the latent variable z
- ③ NN implementation of the encoder gives (for every input x) a vector mean and a diagonal covariance

VAE Decoder



- 1 Decoder takes the latent vector z and returns the parameters for a distribution

VAE Decoder



- ① Decoder takes the latent vector z and returns the parameters for a distribution
- ② $p_{\phi}(x/z)$ gives mean and variance for each pixel in the output

VAE Decoder



- ① Decoder takes the latent vector z and returns the parameters for a distribution
- ② $p_{\phi}(x/z)$ gives mean and variance for each pixel in the output
- ③ Reconstruction of x is via sampling

VAE loss function



- ① Loss for AE: l_2 distance between the input and its reconstruction

VAE loss function



- ① Loss for AE: l_2 distance between the input and its reconstruction
- ② In case of VAE: we need to learn parameters of two probability distributions

VAE loss function



- ① Loss for AE: l_2 distance between the input and its reconstruction
- ② In case of VAE: we need to learn parameters of two probability distributions
- ③ For each input x_i we maximize expected value of returning x_i (or, minimize the NLL)

$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

VAE loss function



$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

- ① Problem: Input images may be memorized in the latent space \rightarrow similar inputs may get different representations in z space

VAE loss function



$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

- ① Problem: Input images may be memorized in the latent space \rightarrow similar inputs may get different representations in z space
- ② We prefer continuous latent representations to give meaningful parameterization (e.g. smooth transition between digits)

VAE loss function



$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

- ① Problem: Input images may be memorized in the latent space \rightarrow similar inputs may get different representations in z space
- ② We prefer continuous latent representations to give meaningful parameterization (e.g. smooth transition between digits)
- ③ Solution: Force $q_{\theta}(z/x_i)$ to be close to a standard distribution (e.g. Gaussian)

VAE loss function



$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)}[\log p_\phi(x_i/z)] + \mathbb{KL}(q_\theta(z/x_i) || p(z))$$


- ① First term promotes recovery, second term keeps encoding continuous (beats memorization)

VAE loss function



$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)}[\log p_\phi(x_i/z)] + \mathbb{KL}(q_\theta(z/x_i) || p(z))$$

① Problem: Differentiating over θ and ϕ

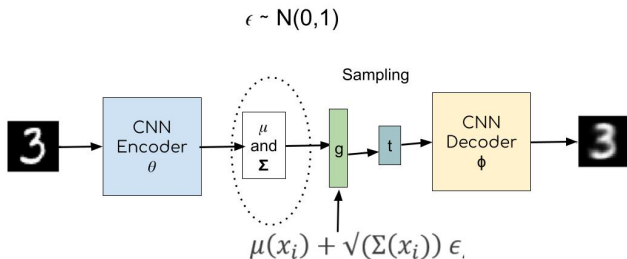


VAE loss function



$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)} [\log p_\phi(x_i/z)] + \text{KL}(q_\theta(z/x_i) || p(z))$$

- ① Reparameterization: Draw samples from $N(0,1) \rightarrow$ doesn't depend on parameters



Generation with VAE



- Sample z from the prior $p(z)$

Generation with VAE



- Sample z from the prior $p(z)$
- Run z through the decoder (ϕ) \rightarrow distribution over data

Generation with VAE



- Sample z from the prior $p(z)$
- Run z through the decoder (ϕ) \rightarrow distribution over data
- Sample from that distribution to generate the sample x

Generation with VAE



Figure credits: Wojceich

Generation with VAE




Figure credits: Kingma et al.

Edit/Manipulate samples with VAE



Slide Title



① Slide content